

# Meta-Programming and Hybrid Parallel Strategies for Solving PDEs: An FDM and PINN Comparison<sup>1, 2</sup>

## Seminar Presentation III

LI YIHAI

Student ID: 23345919

Supervision: Michael Peardon

September 26, 2024



Trinity College Dublin  
Coláiste na Tríonóide, Baile Átha Cliath

<sup>1</sup>Full docs: <https://liyihai.com/html/index.html>

<sup>2</sup>Git repository: <https://github.com/liyihai-official/Final-Project>

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- PINN Model

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- PINN Model

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

# Introduction

## Recap - Deep Neural Network

### Kolmogorov PDEs

Solving  $u(x, T)$ , for  $\mathbb{R}^1 \ni T > 0, x \in \mathbb{R}^d, t \in [0, T], u(t, x) = u \in \mathbb{R}^1, \mu(x) \in \mathbb{R}^d, \sigma(x) \in \mathbb{R}^{d \times d}$ ,

$$u_t = \frac{1}{2} \operatorname{Trace}_{\mathbb{R}^d} [\sigma(x) [\sigma(x)]^* \operatorname{Hess}_x u] + \langle \mu(x), \nabla_x u \rangle_{\mathbb{R}^d} \quad (1)$$

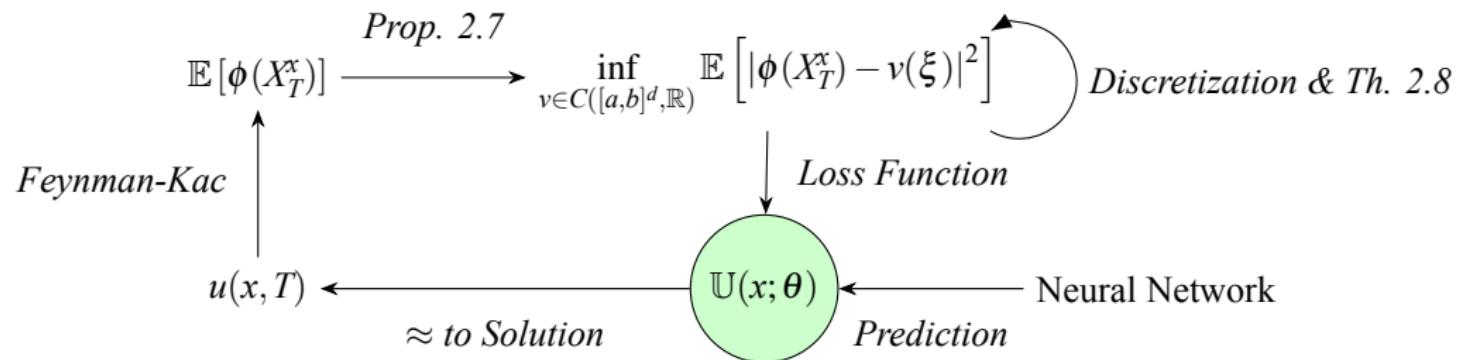


Figure: Deep Neural Network (DNN) Methodology of Solving Kolmogorov PDEs [FIRST]

# Introduction

## Recap - Physics Informed Nerual Network

### General Form of PDEs

-  $u(t,x)$  denotes with the target function,  $x \in \mathbb{R}^d$ .

-  $\Gamma [\cdot ; \lambda]$  is a non-linear operator parameterized by  $\lambda$ .

$$u_t(t,x) + \Gamma [u; \lambda] = 0 \quad (2)$$

Define  $f(t,x)$  to be given by

$$f(t,x) = u_t(t,x) + \Gamma [u; \lambda] \quad (3)$$

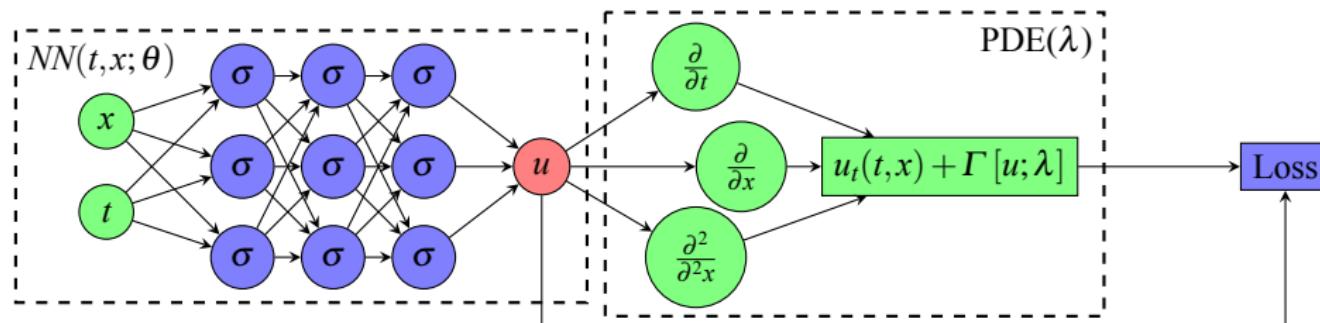


Figure: PINN, with 3 fully connected hidden layers

### Comparing With Finite Difference Time Domain Method (FDTD)

#### ① Deep Neural Network [FIRST]

- Gives lower quality approximations.
- Takes longer time to train.
- Possible to solve high dimension PDEs

#### ② Physics Informed Neural Network

- Gives higher quality approximations.
- Takes longer time to train.
- Has more flexible way to get results.
- Possible to solve high dimension PDEs

# Objectives

This project focused on following objectives:

- ① Implement FDTD and PINN in [C++/C](#).
- ② Implement FDTD hybrid parallel version using [MPI/OpenMP](#).
- ③ Implement PINN GPU parallel version using [Libtorch/CUDA](#).
- ④ Evaluate the efficiency and accuracy of FDTDs and PINNs.

# Challanges

However, there were many obstacles including:

- ① Portability, Continuity.
- ② Overlapping Communication and Computation.
- ③ Unnecessary intra-node communication
- ④ Communication Overhead.
- ⑤ Scalability Issues.
- ⑥ Memory Management.

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- PINN Model

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

# General Form of problem

The PDE parametrized by number  $\lambda$  and an operator  $\mathcal{N}[\cdot; \lambda]$ , and assume the variable  $x$  is a 2D or 3D spatio-vector which is written in

$$\begin{cases} \frac{\partial u}{\partial t}(t, \vec{x}) + \mathcal{N}[u; \lambda] = 0 \\ u(0, \vec{x}) = \varphi(\vec{x}) \end{cases} \quad (4)$$

where  $\varphi$  is the initial condition, and  $\vec{x} \in \Omega, t \in [0, +\infty)$ .

## Boundary Conditions

The Dirichlet and Von Neumann boundary conditions are formed as

$$\begin{cases} u(t, \vec{x}) = g(t, \vec{x}) \\ \frac{\partial u}{\partial \vec{n}} = g(t, \vec{x}) \end{cases} \quad (5)$$

where  $\vec{n}$  is the normal vector on  $\overline{\Omega}$  the boundary of domain  $\Omega$ .

# Thermal Conduction Systems

## Heat Equation 2D

The function

$$u(t,x,y) = x + y - xy, \forall \alpha \in \mathbb{R}^1 \quad (6)$$

is the solution of 2D Heat Equation 7 below

$$\frac{\partial u}{\partial t} = \alpha \left( \frac{\partial u^2}{\partial^2 x} + \frac{\partial u^2}{\partial^2 y} \right) \quad (x,y) \in \Omega, t \in [0, +\infty)$$
$$u(0,x,y) = \varphi(x,y) = 0 \quad (x,y) \in \Omega \quad (7)$$

$$u(t,x,y) = g(x,y) = \begin{cases} y, & x = 0, y \in (0,1) \\ 1, & x = 1, y \in (0,1) \\ x, & y = 0, x \in (0,1) \\ 1, & y = 1, x \in (0,1) \end{cases} \quad t \in [0, +\infty)$$

# Thermal Conduction Systems

## Heat Equation 3D

The function

$$u(t, x, y, z) = x + y + z - 2xy - 2xz - 2yz + 4xyz, \forall \alpha \in \mathbb{R}^1 \quad (8)$$

is the solution of 3D Heat Equation 9 below

$$\begin{aligned} \frac{\partial u}{\partial t} &= \alpha \left( \frac{\partial u^2}{\partial^2 x} + \frac{\partial u^2}{\partial^2 y} + \frac{\partial u^2}{\partial^2 z} \right) & (x, y, z) \in \Omega, t \in [0, +\infty) \\ u(0, x, y, z) &= \varphi(x, y, z) = 0 & (x, y, z) \in \Omega \end{aligned} \quad (9)$$

$$u(t, x, y, z) = g(x, y, z) = \begin{cases} y + z - 2yz, & x = 0, \\ 1 - y - z + 2yz, & x = 1, \\ x + z - 2xz, & y = 0, \\ 1 - x - z + 2xz, & y = 1, \\ x + y - 2xy, & z = 0, \\ 1 - x - y + 2xy, & z = 1 \end{cases} \quad t \in [0, +\infty)$$

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- PINN Model

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

# N-dimension Matrix

## Challanges

STL provides containers `std::array` and `std::vector` for creating one-dimension array.  
There are two way for dealing with high-dimension data:

- ① Nesting the one dimension arrays or vectors.
- ② Hierarchy approach, designing derived classes of one-dimension array base class.

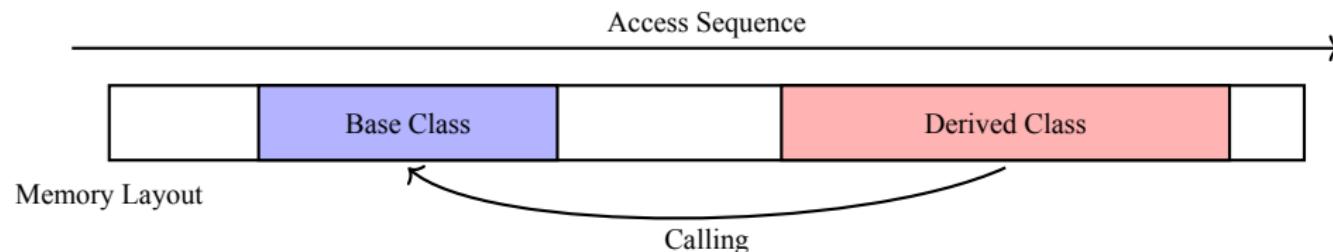


Figure: Derived Class calling members in Base class, timing is not predictable.

- ① Nesting multi-dimension array has non-contiguous memory layout.
- ② Derived class needs more time to access members in base class.
- ③ Poor cache utilization leads to poor performance.
- ④ MPI type create requires contiguous memory layout.

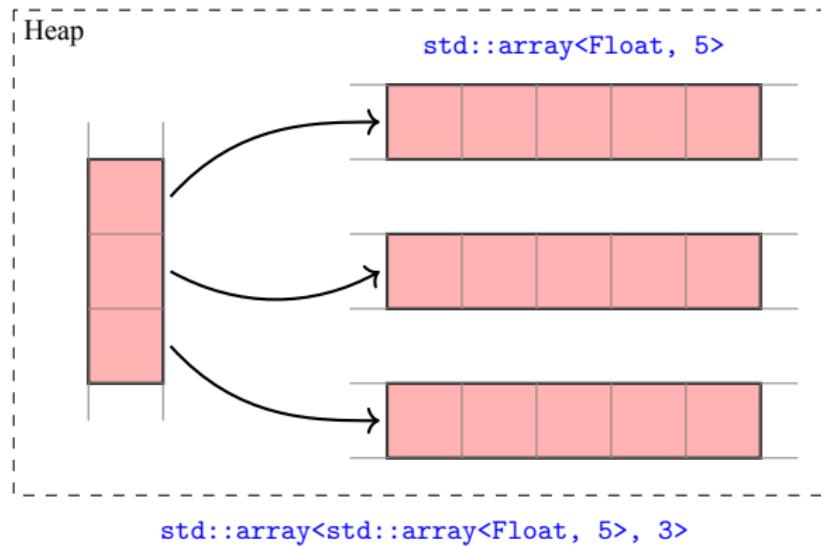


Figure: Using nested `std::array<T, N>` to store 2D array data.

# N-dimension Matrix

## Solution

- Separate into two detail and user interface objects adhering RAII rules.

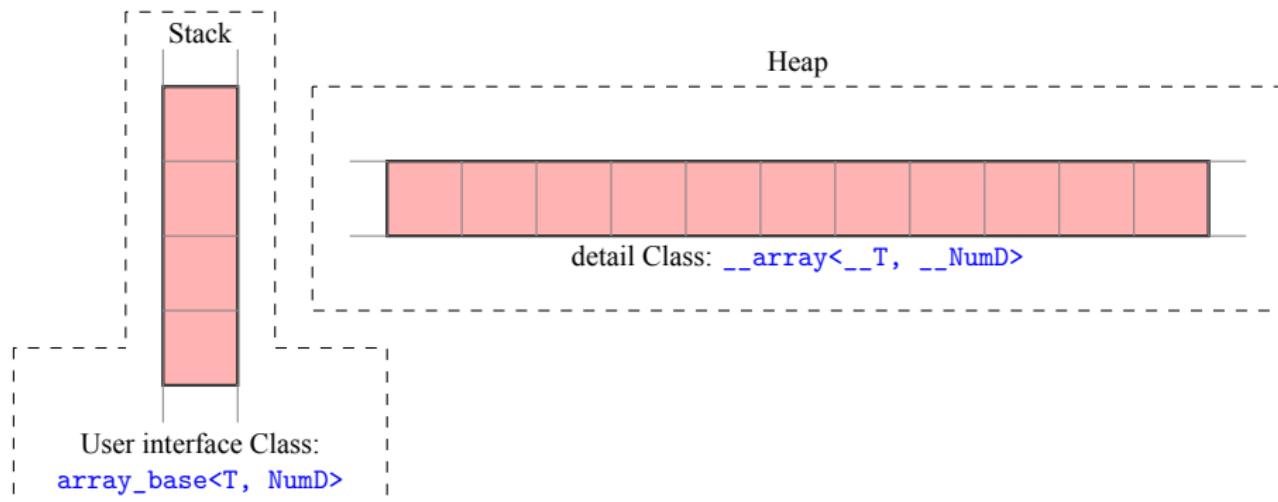


Figure: The solution of N-dimension Matrix, using detail Class, user interface Class

# N-dimension Matrix

## Solution

- Separate into two detail and user interface objects adhering RAII rules.
- An external small `_multi_array_shape` object defines the routines for accessing the elements.

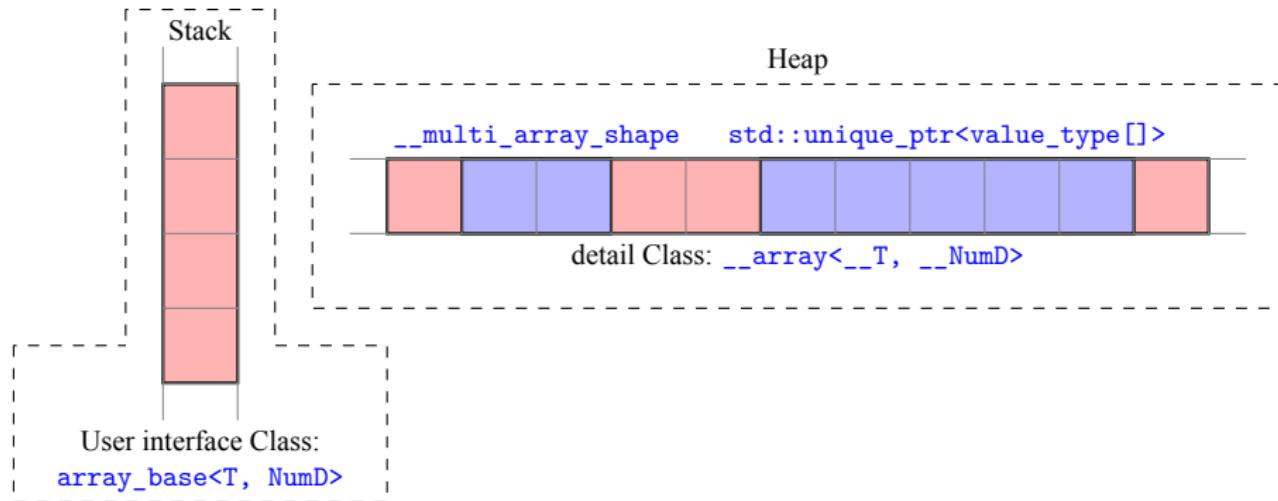


Figure: The solution of N-dimension Matrix, using detail Class, user interface Class and a shape management structure.

# N-dimension Matrix

## Solution

- Separate into two detail and user interface objects adhering RAII rules.
- An external small `_multi_array_shape` object defines the routines for accessing the elements.
- Smart pointer, ensure memory's contiguous layout and safety.

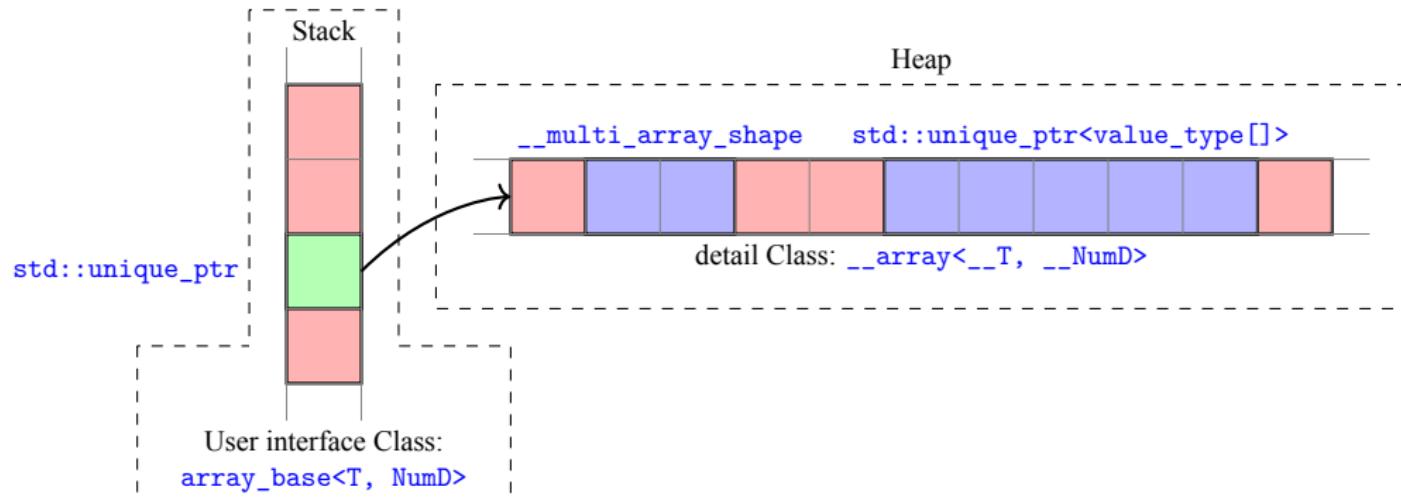


Figure: The solution of N-dimension Matrix, using detail Class, user interface Class and a shape management structure.

# Parallelization of N-dimension Arrays

## MPI Environment

The hybrid PDE solver requires the MPI supports multi-threads on each processes.

- High-level libraries like Boost.MPI
  - ① have better MPI resource management and other basic communication features.
  - ② only provide limited useful features for latter PDE solvers.
  - ③ lead to lower performance than low-level OpenMPI.

# Parallelization of N-dimension Arrays

## MPI Environment

The hybrid PDE solver requires the MPI supports multi-threads on each processes.

- High-level libraries like Boost.MPI
  - ① have better MPI resource management and other basic communication features.
  - ② only provide limited useful features for latter PDE solvers.
  - ③ lead to lower performance than low-level OpenMPI.
- I developed an environment class for MPI
  - ① better resource management than raw MPI.
  - ② provides basic features exclusively for this project.

# Parallelization of N-dimension Arrays

## MPI Topology - Challenges

Distributed N-dimension arrays are created based on MPI N-dimension Cartesian topology.

# Parallelization of N-dimension Arrays

## MPI Topology - Challenges

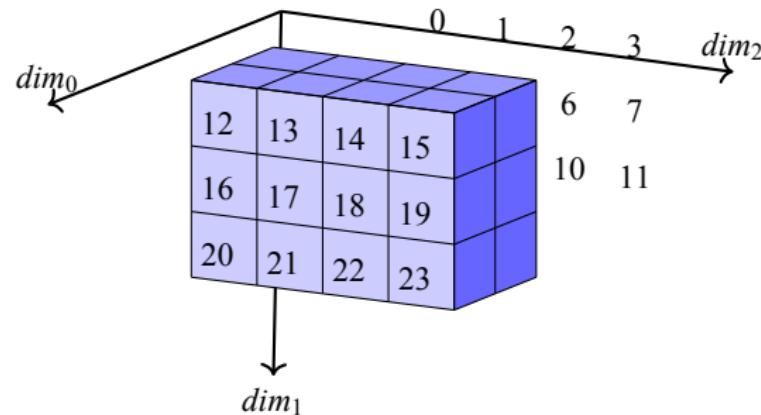
Distributed N-dimension arrays are created based on MPI N-dimension Cartesian topology.

- ① Ghost communication is required in overlapping of MPI communication and local computation.
- ② Parallel I/O is needed for debugging and storing results.
- ③ Topology information will be frequently used in PDE solver.

# Parallelization of N-dimension Arrays

## MPI Topology - Solution

- ① An MPI Topology class defines the distribution details of N-dimension array.



**Figure:** The Schematic representation of the MPI Cartesian topology scheme of 24 processors on 3 dimension space, which has a  $2 \times 3 \times 4$  grid of processes

# Parallelization of N-dimension Arrays

## MPI Topology - Solution

- ② Using `MPI_Type_create_subarray` for creating Ghost MPI datatype for communication.

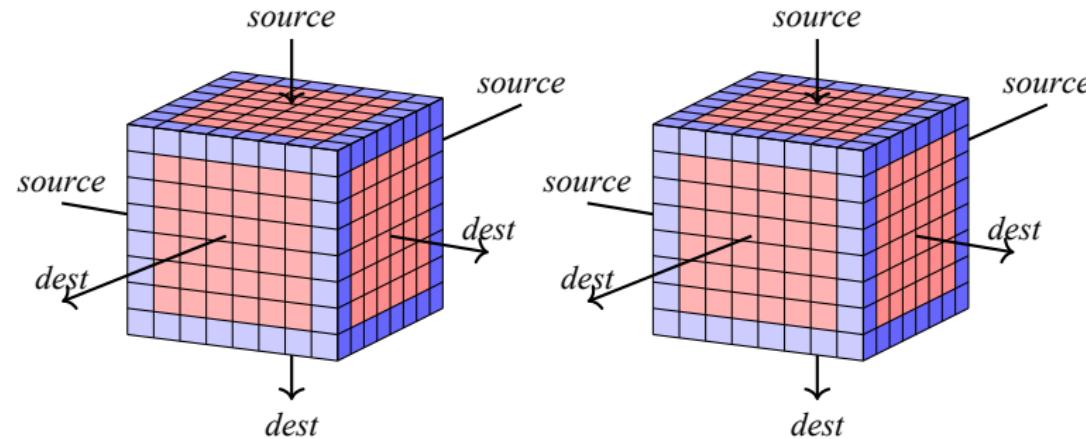


Figure: Representation of a 3D MPI Communication Scheme of 2 among 24 processes between 3 dimension sub-arrays.

# Parallelization of N-dimension Arrays

## MPI Topology - Solution

- ③ Cartesian array has members Topology class and N-dimension array class, to ensure they are closely located on memory.
- ④ The external functions provide gather-based I/O and MPI I/O for handling different scenarios.
  - gather-based function is designed for small scale problem,  
it collects all data included with the boundaries to the root process, and use the I/O of N-dimension array.
  - MPI I/O function is designed for handling large scale problem,  
it only collects the internal bulk without boundary.

# Parallelization of N-dimension Arrays

## MPI Topology - Solution

- ⑤ User interface class unified features for both topology and array classes.
- ⑥ Separate into two detail and user interface objects adhering RAII rules.

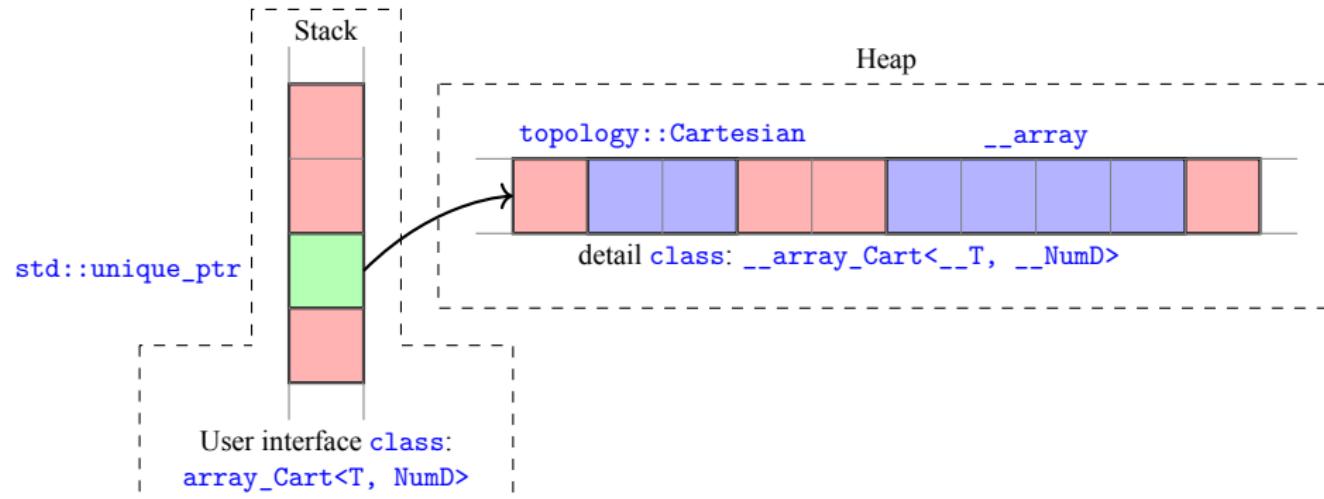


Figure: The solution of distributed N-dimension Array, using detail `class`, user interface `class` and a Cartesian topology `class`.

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- PINN Model

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

## Implementations

- PDE Solvers
  - Boundary/Initial Conditions
  - Pure MPI
  - Master-only, no overlapping
  - Master-only, with overlapping
- PINN Model

# N-dimension Boundary and Initial Conditions

## Challanges

- ① Mathematical function needs be discretized and distributed as well.
- ② Initial and Dirichlet boundary conditions only applies once.
- ③ Von Neumann boundary condition participates the evolving process in FDM.
- ④ Need to access the data.

# N-dimension Boundary and Initial Conditions

## Solution

- ① Use `lambda` function to construct classes.
- ② Creating external classes for each conditions as the `friend` classes of PDE solver classes.
- ③ Set `Bool` vectors help to determine the status of conditions and type of boundary conditions.

# N-dimension PDE solvers

## Challanges

- ① PDE solver of Heat Equations in different dimension space have similar parameters and features, but type-field solution lead to code redundancy.
- ② Applying MPI communications between local arrays, avoiding overhead.

# N-dimension PDE solvers

## Challanges

- ① PDE solver of Heat Equations in different dimension space have similar parameters and features, but type-field solution lead to code redundancy.
- ② Applying MPI communications between local arrays, avoiding overhead.
- ③ Three type of strategies:
  - ① Pure MPI parallel
  - ② Master-only, no overlapping hybrid parallel.
  - ③ Master-only, communication/computation (comm./comp.) overlapping hybrid parallel.

# N-dimension PDE solvers

## Solutions

`virtual` function is resolved at run-time through `vtable`, and only lose up to about 25% efficiency in terms of the function call mechanism,

- ① Create an abstract Heat base class of Heat Equation, and `override` functions in derived class on every dimension.

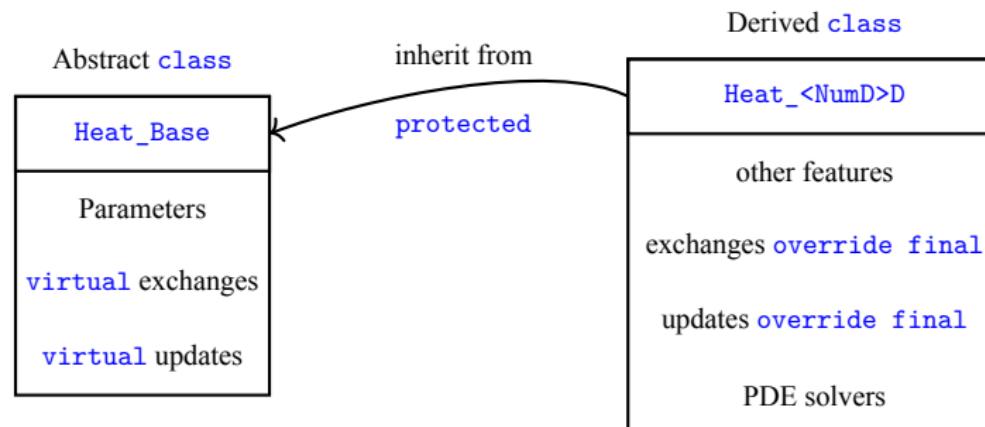


Figure: Abstract inheritance scheme representation.

# N-dimension PDE solvers

## Solutions

The MPI Remote Memory Access (MPI-RMA) operations require hardware support to gain the benefits, so:

- ② MPI communications are implemented in
  - non-blocking using `MPI_Isend/MPI_Irecv`.
  - blocking using `MPI_Sendrecv`

## Implementations

- PDE Solvers
  - Boundary/Initial Conditions
- Pure MPI
  - Master-only, no overlapping
  - Master-only, with overlapping
- PINN Model

# N-dimension PDE solvers

## Solutions

- ③ Three type of strategies
  - ① Pure MPI parallel

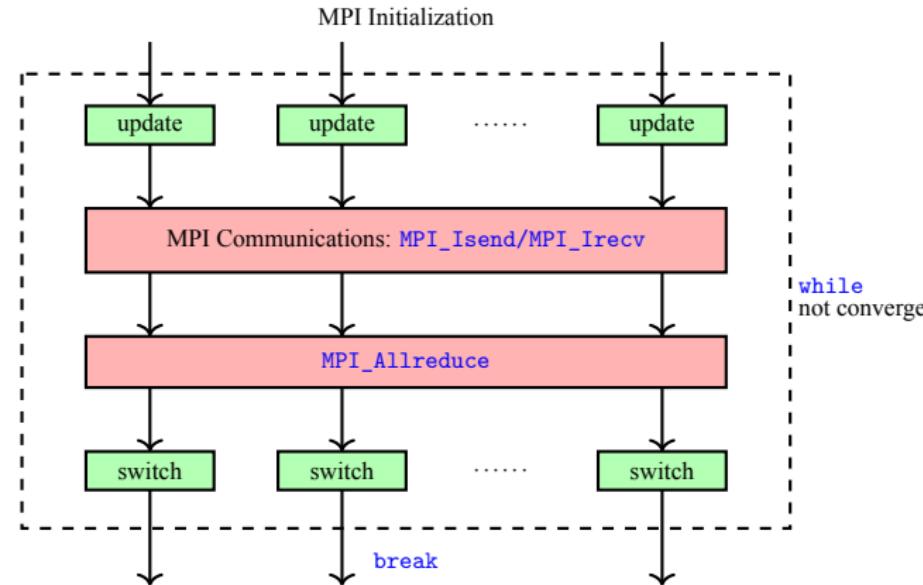


Figure: Scheme Representation: flow chart of Pure MPI model

# N-dimension PDE solvers

## Solutions

### ③ Three type of strategies

- ① Master-only, no overlapping hybrid parallel.

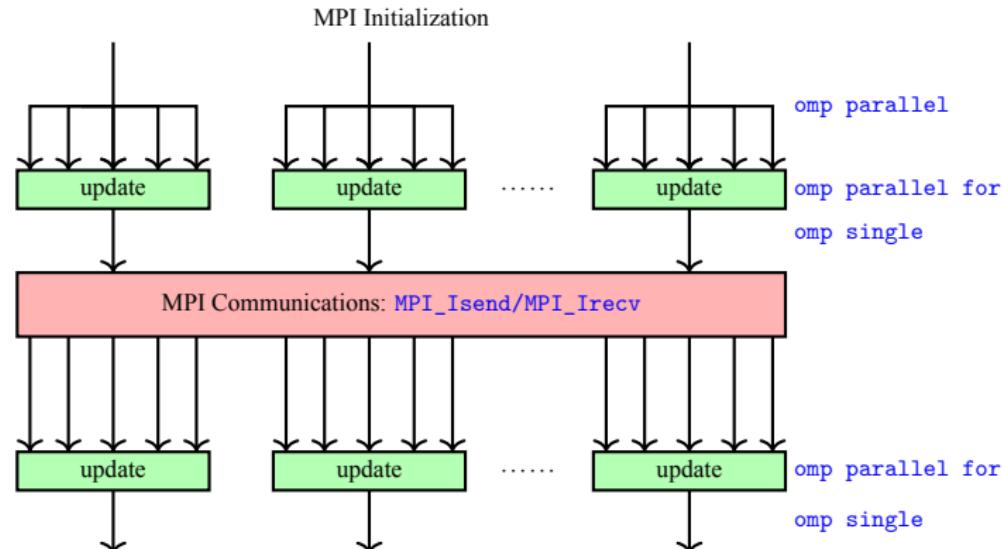


Figure: Scheme Representation: flow chart of Hybrid model, with no overlapping

# N-dimension PDE solvers

## Solutions

### ③ Three type of strategies

- ① Master-only, comm./comp. overlapping hybrid parallel.

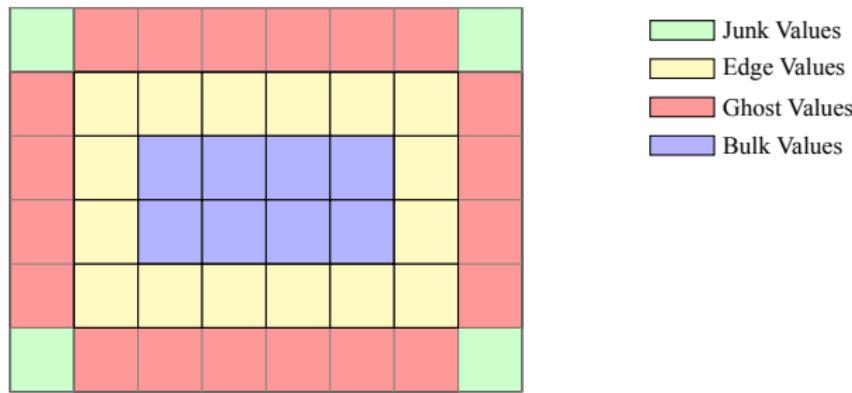


Figure: Scheme Representation: Grid 2D, type of values

# N-dimension PDE solvers

## Solutions

### ③ Three type of strategies

- ① Master-only, comm./comp. overlapping hybrid parallel.

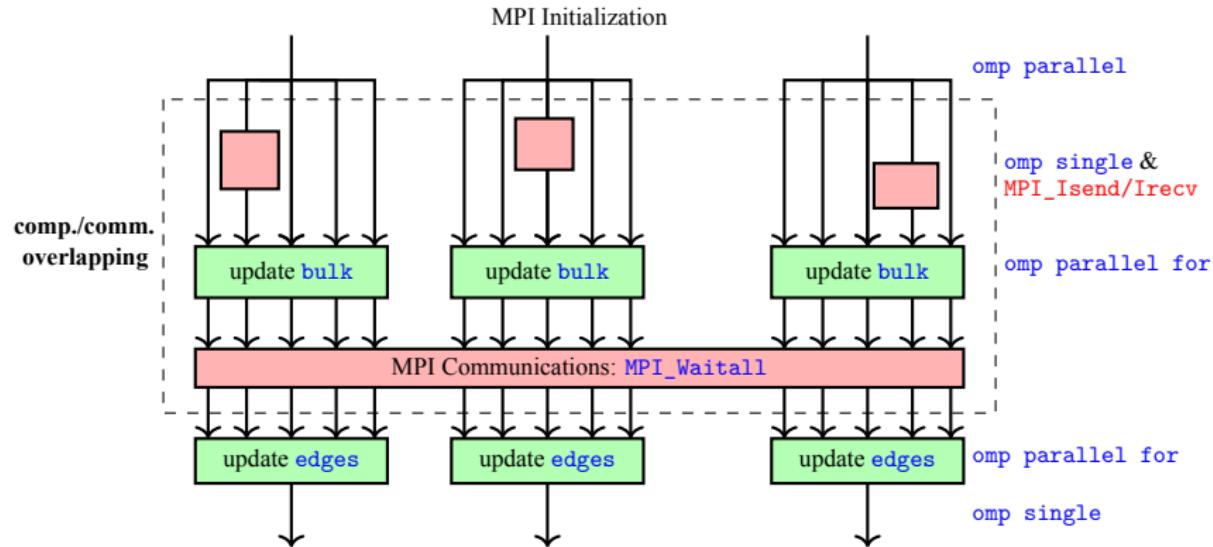


Figure: Scheme Representation: flow chart of Hybrid model, with comp./comm. overlapping

# PINN Model

## Challanges

Python has many easy-use libraries such as Pytorch, Tensorflow and Caffe.

# PINN Model

## Challanges

Python has many easy-use libraries such as Pytorch, Tensorflow and Caffe.  
For neural network implementations, there are some issues in practice:

- ① Interpret language Python is significant slower than compile language C/C++
- ② Worse resource management than C/C++
- ③ Lower security in parallel program.

# PINN Model Solution

For getting higher performance and better safety, I chose to use

- ➊ Pytorch C++ API: [Libtorch](#).

to reproduce the Python version of PINN.

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- Boundary/Initial Conditions

- Pure MPI
- Master-only, no overlapping
- Master-only, with overlapping

## • PINN Model

- Customized Loss Function
- Network Structure

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

## GPU cluster

Each node has a

- AMD Ryzen 7, 7800 X3D, with 16 CPU(s)
- NVIDIA RTX 4090, 24GB

# Hardware

## CPU cluster

Each node has a: Intel(R) Xeon(R) Platinum 9242 CPU 2.30Ghz  
with 96 CPU(s), 4 Non-unified Memory Access (NUMA) nodes, 1MB L2 cache, 35.75MB L3 cache.

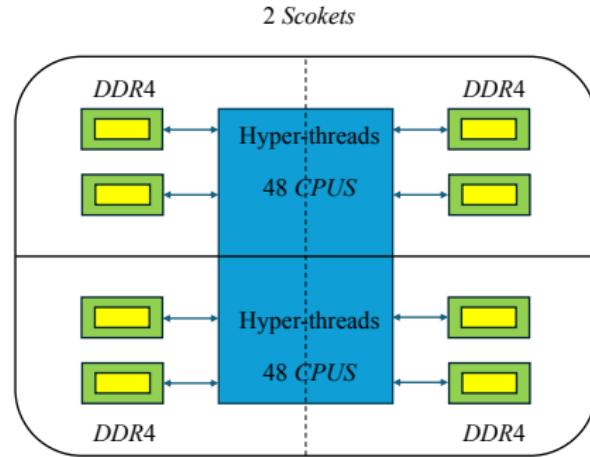


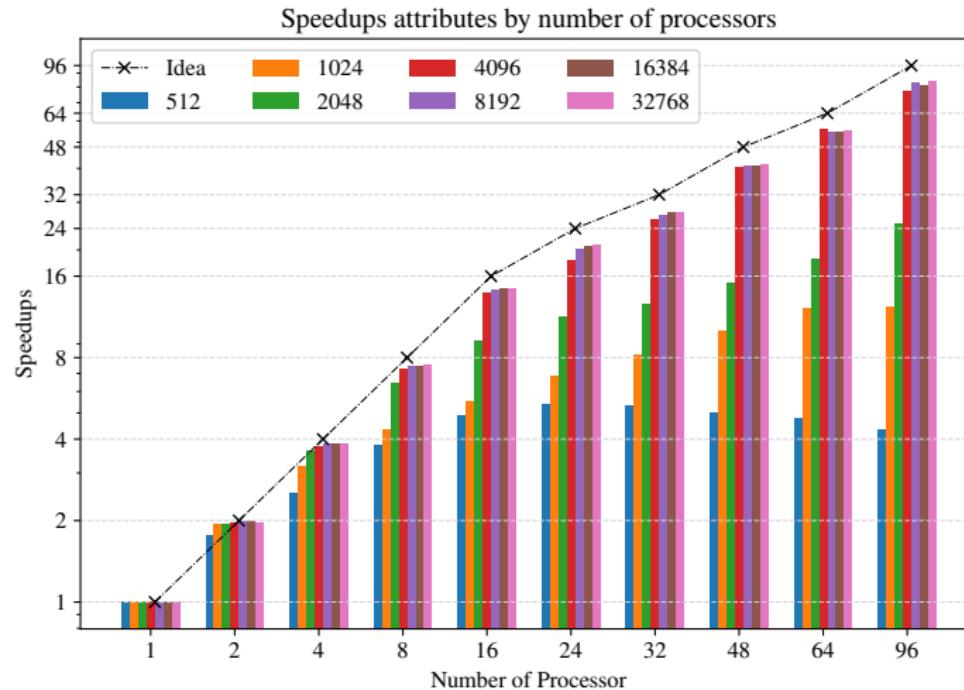
Figure: NUMA topology of single node on Cluster

# Single Node Scaling Tests

- 2D Heat Equation
- Single Precision
- Problem Size:  $512^2, 1024^2, 2048^2, 4096^2, 8192^2, 16384^2, 32768^2$
- Run 5 times for each test, remove the max/minum values, use the average of left 3.

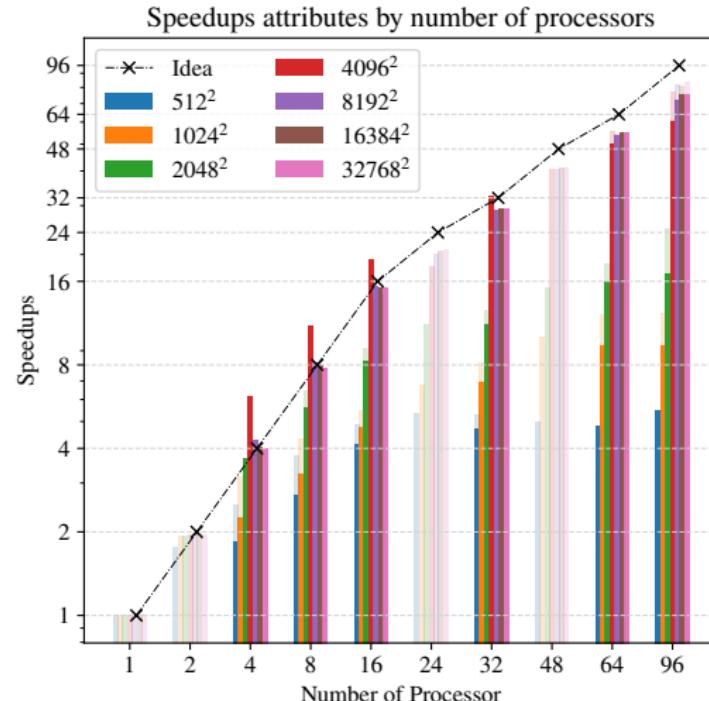
# Single Node Scaling Tests

## Strong Scaling - Pure MPI

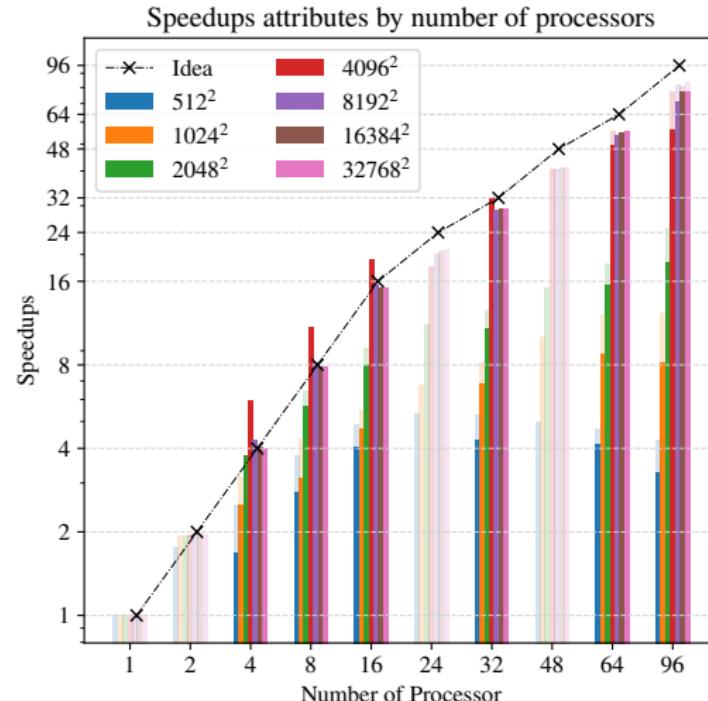


# Single Node Scaling Tests

## Strong Scaling - MPI/OpenMP Hybrid



(a) No overlapping comm./comp.



(b) With overlapping comm./comp.

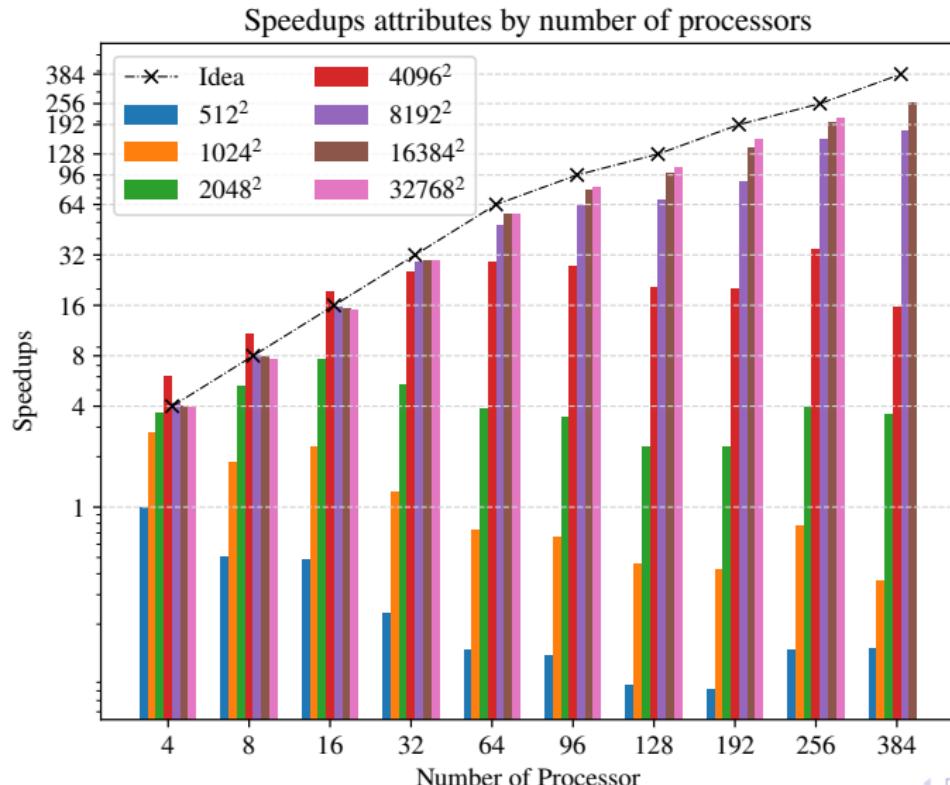
# Single Node Scaling Tests

## Weak Scaling

Strategy	Size	Number of CPUs			$f_p(\%)$
		4	16	64	
<b>Pure MPI</b>	$512^2$	4.006	12.497	47.849	75.1
<b>No Overlap</b>		2.876	11.206	42.754	67.0
<b>With Overlap</b>		3.173	10.818	42.282	66.2
<b>Pure MPI</b>	$1024^2$	3.838	9.304	33.707	53.2
<b>No Overlap</b>		3.947	12.995	33.447	54.1
<b>With Overlap</b>		4.024	12.932	33.361	54.0
<b>Pure MPI</b>	$2048^2$	2.376	8.245	31.203	49.0
<b>No Overlap</b>		3.874	8.972	31.510	49.8
<b>With Overlap</b>		3.740	8.989	31.430	49.7
<b>Pure MPI</b>	$4096^2$	3.543	8.245	31.203	77.5
<b>No Overlap</b>		3.953	13.799	49.515	78.0
<b>With Overlap</b>		3.948	13.800	49.989	78.7

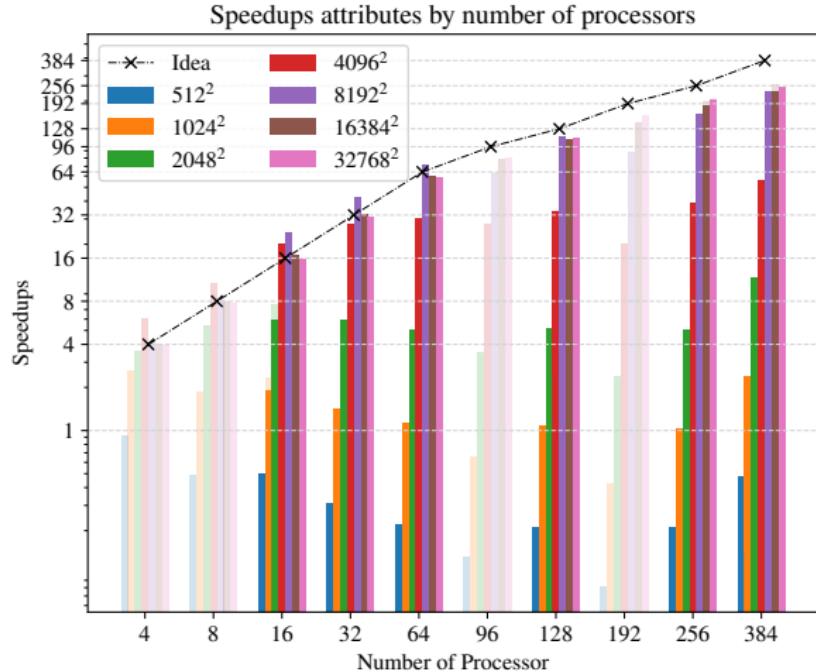
# Multi-node Scaling Tests

## Strong Scaling - Pure MPI

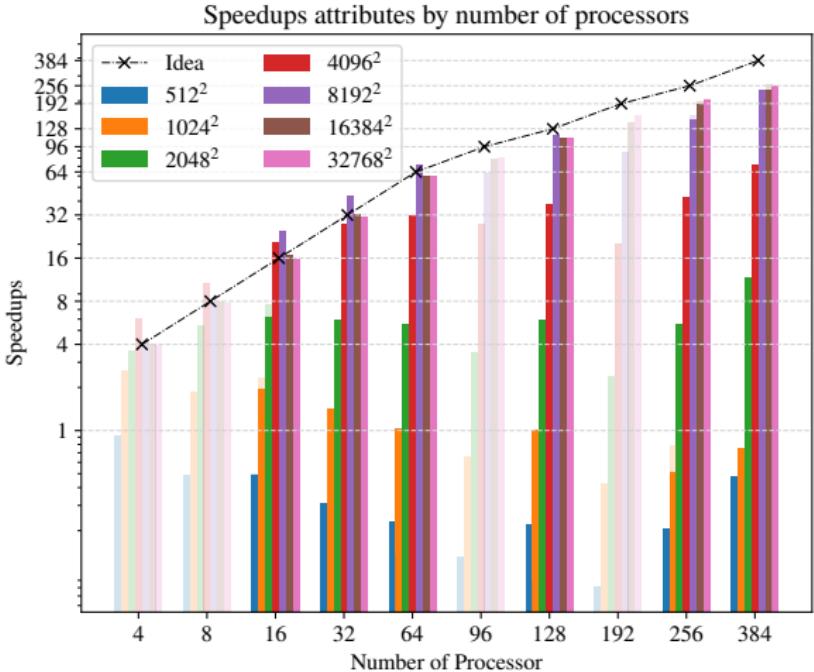


# Multi-node Scaling Tests

## Strong Scaling - MPI/OpenMP Hybrid



(a) No overlapping comm./comp.



(b) With overlapping comm./comp.

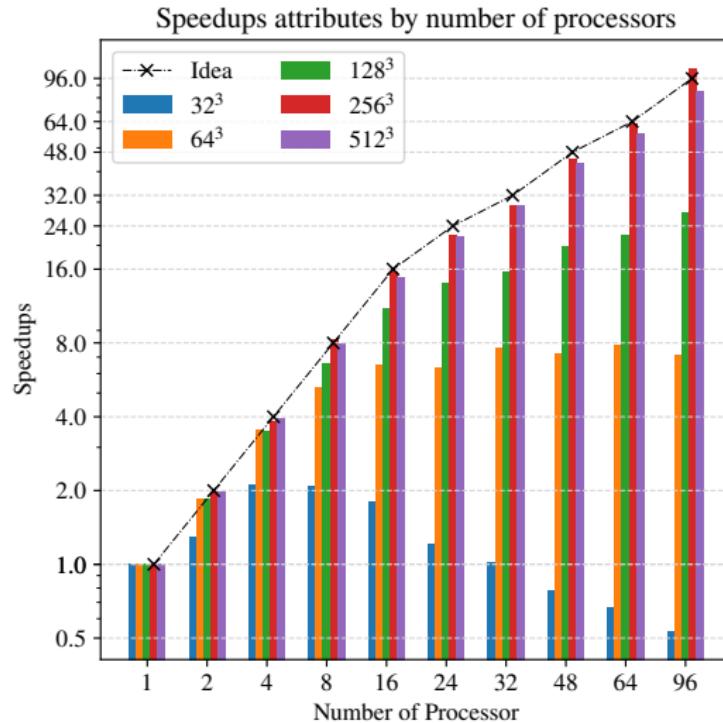
# Multi-node Scaling Tests

## Weak Scaling

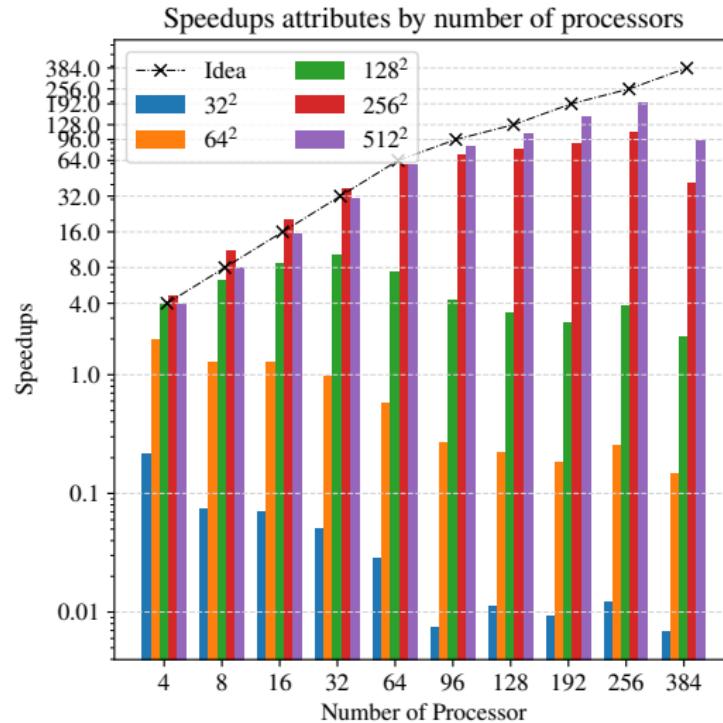
Strategy	Size	Number of CPUs				$f_p(\%)$
		4	16	64	256	
<b>Pure MPI</b>	512 <sup>2</sup>	3.300	10.379	26.000	124.375	48.2
<b>No Overlap</b>		-	8.094	25.931	127.648	49.3
<b>With Overlap</b>		-	8.386	27.126	116.951	45.5
<b>Pure MPI</b>	1024 <sup>2</sup>	3.848	13.037	30.143	-	49.3
<b>No Overlap</b>		-	13.702	44.040	-	69.8
<b>With Overlap</b>		-	13.834	44.090	-	69.9
<b>Pure MPI</b>	2048 <sup>2</sup>	3.787	9.172	32.013	-	50.6
<b>No Overlap</b>		-	13.936	34.368	-	55.7
<b>With Overlap</b>		-	14.274	34.414	-	55.9
<b>Pure MPI</b>	4096 <sup>2</sup>	3.884	13.859	51.041	-	80.2
<b>No Overlap</b>		-	15.315	53.157	-	83.8
<b>With Overlap</b>		-	15.294	53.401	-	84.2

# 3D Space Heat Equation

## Strong Scaling - pure MPI



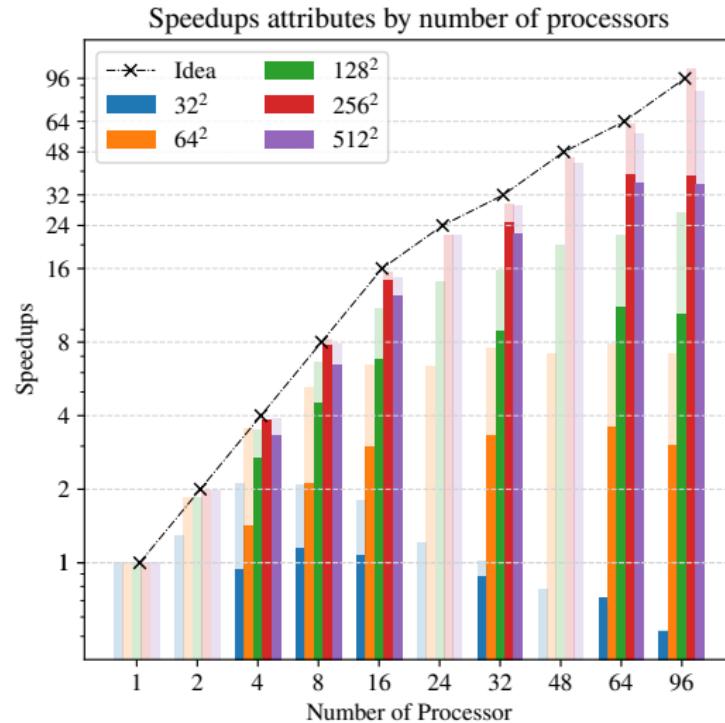
(a) Single node



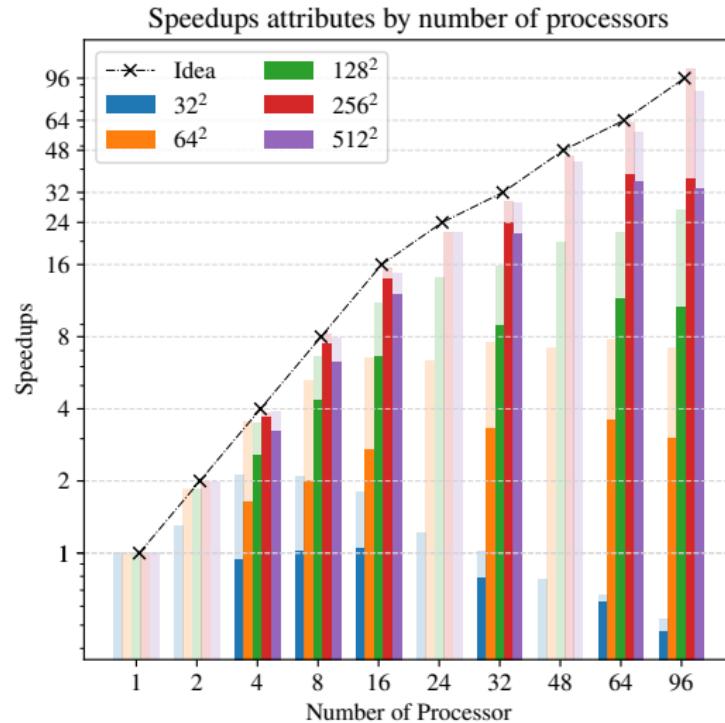
(b) Four nodes

# 3D Space Heat Equation

Strong Scaling - MPI/OpenMP Hybrid, Single Node



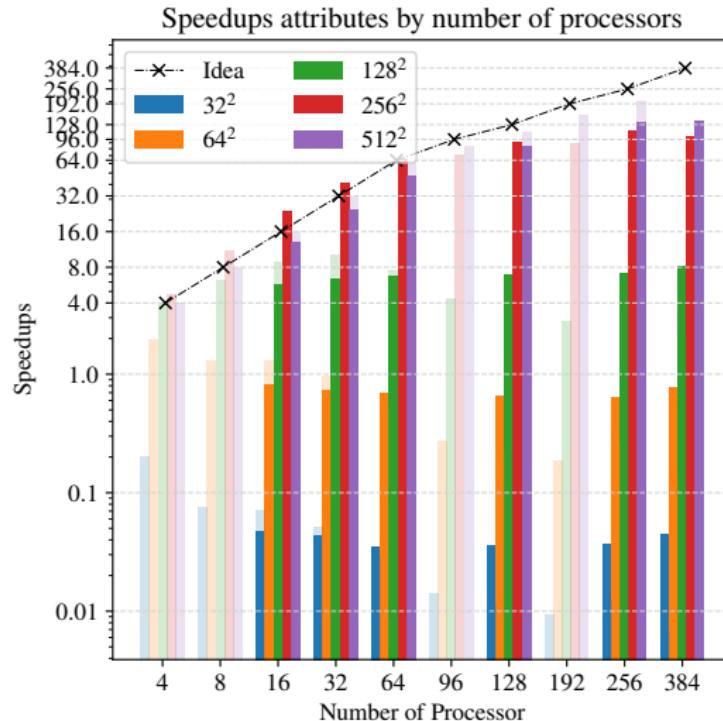
(a) Non-overlapping



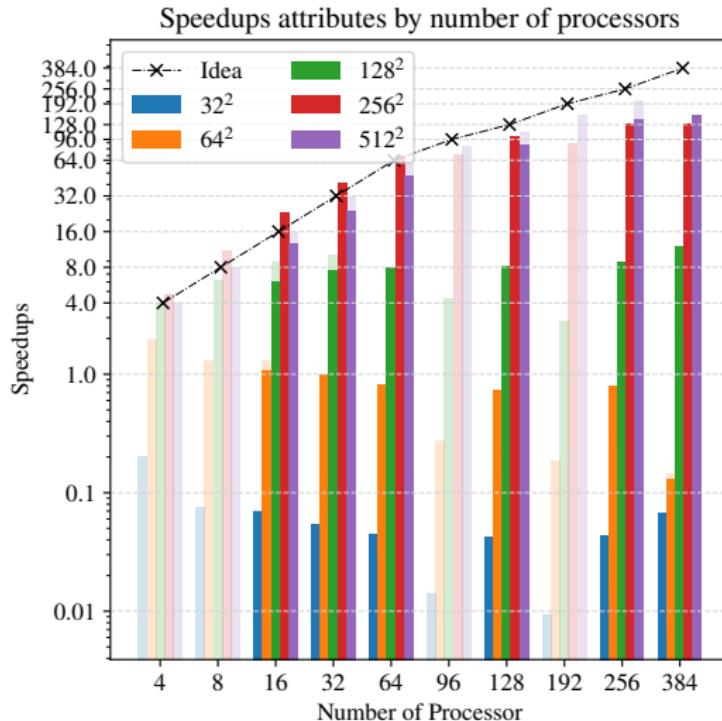
(b) With overlapping

# 3D Space Heat Equation

## Strong Scaling - MPI/OpenMP Hybrid, Four Nodes



(a) Non-overlapping



(b) With overlapping

# 3D Space Heat Equation

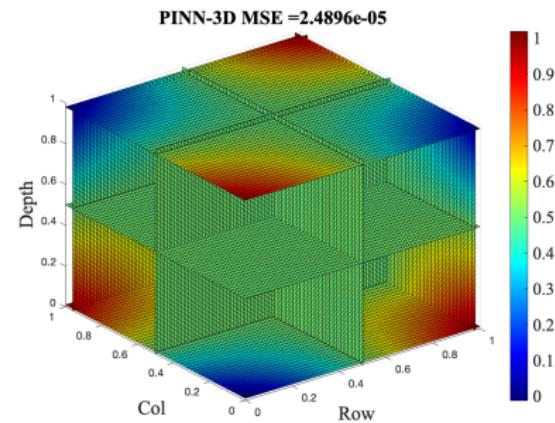
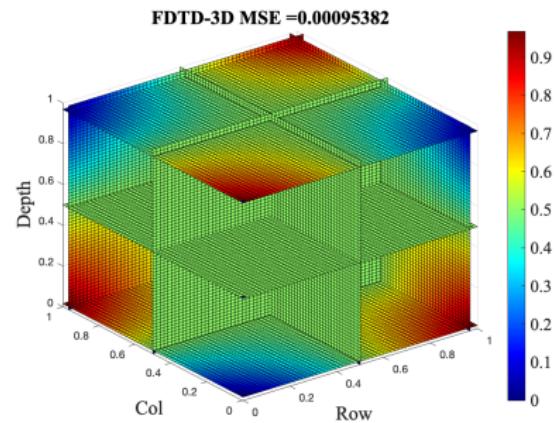
## Weak Scaling

Strategy	Size	Number of CPUs			$f_p(\%)$	$f_p(\%)$ Multi-node
		8	64	64 Multi-node		
<b>Pure MPI</b>	$32^3$	5.243	26.405	9.078	41.6	14.2
<b>No Overlap</b>		2.124	13.386	8.094	21.0	12.7
<b>With Overlap</b>		2.014	13.884	9.586	21.8	39.7
<b>Pure MPI</b>	$64^3$	7.937	32.034	30.106	50.8	47.1
<b>No Overlap</b>		5.393	20.007	33.460	31.8	52.3
<b>With Overlap</b>		5.200	19.558	35.254	31.1	31.6
<b>Pure MPI</b>	$128^3$	3.513	24.239	25.428	38.0	39.7
<b>No Overlap</b>		3.303	15.235	35.254	24.1	55.1
<b>With Overlap</b>		3.187	15.107	20.096	23.9	31.4

# PINN v.s. FDTD Accuracy

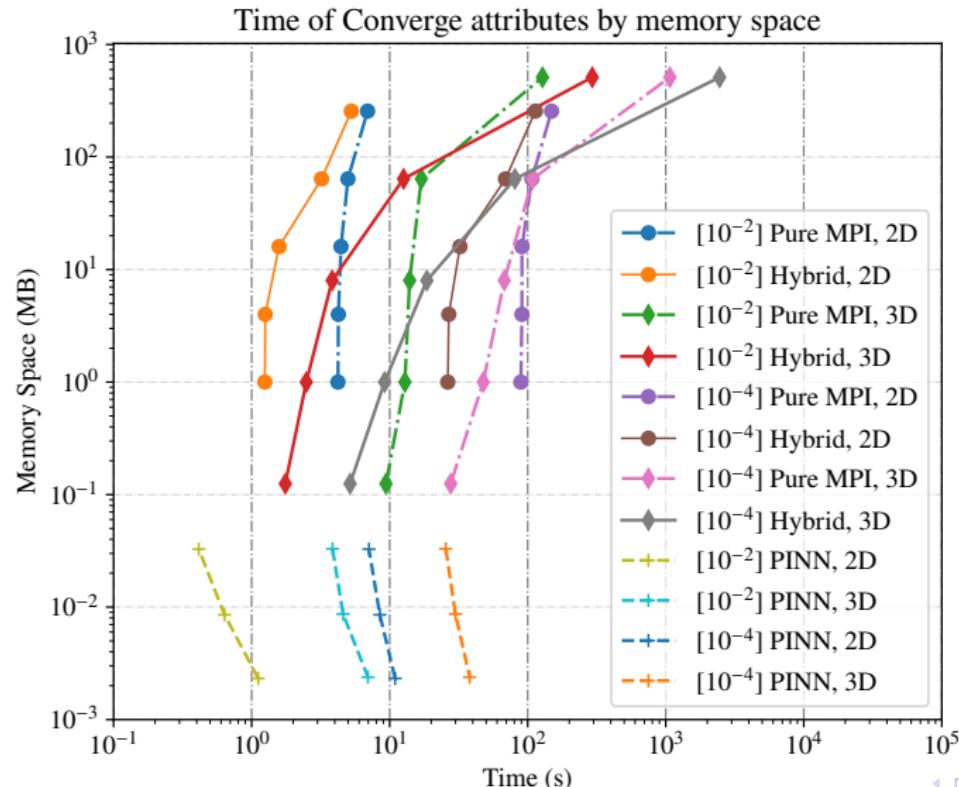
Table: Mean Square Error (MSE) of Approximations, Tolerance:  $10^{-4}$

Method	Heat 2D	Heat 3D
FDTD	1.7287	953.84
PINN	13.37	24.896



# PINN v.s. FDTD

## Memory Usage v.s. Time of Convergence



# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- Boundary/Initial Conditions

- Pure MPI
- Master-only, no overlapping
- Master-only, with overlapping

## 5 PINN Model

- Customized Loss Function
- Network Structure

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

# Conclusion

## Finite Difference Time Domain Method

For 2D and 3D Thermal Conduction PDE systems

I implemented the FDTD methods in fined region with three parallel models:

- Pure MPI model
  - ① had best performance on single compute node overall.
  - ② speedup ratio quickly drop on multi-node.
- MPI/OpenMP Hybrid models
  - ① had lower performance in general.
  - ② can make more use of the advantage of L3 cache.
  - ③ superliner speedup in certain scenario.

# Conclusion

## Physics Informed Neural Network

For 2D and 3D Thermal Conduction PDE systems

I implemented PINN models and trained on generated datasets, comparing with FDTD, the PINN

- ① had identical or higher accuracy.
- ② had less memory usage.
- ③ took shorter time to get the predictions(results).
- ④ had more flexibility on producing the results.

# Conclusion

## Overall Solution

In this project, in order to overcome the obstacles, I

- Used abstract mechanisms in C++, the Meta-programming techniques, to ensure high flexibility, security.
- Used smart pointers to have better resource management.
- Tried to use hybrid parallel techniques to have better scaling performance.
- Made a comparison with the state-of-the-art PINN models.

# Outline

## 1 Introduction

- Related Work
- Challenges & Objectives

## 2 Problem Setups

- General Form
- Thermal Conduction Systems

## 3 Methodology

- N-Dimension Matrix
- Parallelization of N-dimension Arrays

## 4 Implementations

- PDE Solvers
- Boundary/Initial Conditions

- Pure MPI
- Master-only, no overlapping
- Master-only, with overlapping

## • PINN Model

- Customized Loss Function
- Network Structure

## 5 Experiments

- On Single Node
- On Multi-node
- On Different Dimensions
- With PINN on accuracy

## 6 Conclusion

## 7 Further Research Directions

# Further Research Directions

- Resource Management
- Workload Management
- MPI/CUDA Hybrid parallel of FDTD/PINN
- Other PDE Systems
- Other Boundary Conditions

# Closing Remarks <sup>3, 4</sup>

Thank you for your attention!

Any questions?

---

<sup>3</sup>Full docs: <https://liyihai.com/html/index.html>

<sup>4</sup>Git repository: <https://github.com/liyihai-official/Final-Project>